

Analyzing Video Game Sales

Jacob Miller

No need to request edit permission. To make a copy for your presentation, use “File->Download”.

Abstract

The dataset analyzed comes from Kaggle (<https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>) and captures video game sales and Metacritic ratings for video games ranging from 1980 to 2016. By looking at sales based on platform, genre, geographical region, developer and game rating (age appropriateness), I aim to get a good understanding of the optimal markets to target for a video game developer. Action, Shooter and Sports games all found success well above any other genre. The Xbox 360 clearly won as a console last generation, while the PS4 looks to be taking the lead for this generation.

Motivation

I have been a video game player for my entire life - at one time strongly considering going into game development. If I was working at one of the big game development studios, it would be very useful to know what genre of game sold the best, which platform was the most popular, and which geographical region to target. As a generalization, the big studios are more focused on sales numbers than anything else (e.g. creating a brand new game franchise), so having this information will help them decide which markets to pursue. On the other hand, this information might not be very useful to an indie game developer, who is working on their passion project no matter what the sales figures suggest.

Dataset(s)

The dataset used comes from Kaggle (<https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>) and captures data on 16,719 video games from 1980-2016. It contains basic data on the games (platform, publisher, developer, year of release), sales in various geographical regions and game scores based on user and Metacritic reviews. The geographical regions are broken into the largest markets (North America, Europe, Japan), and then the rest of the world.

Data Preparation and Cleaning

This dataset contains a lot of missing (NaN) values. Of the 16,719 rows of data, 9,894 of them have at least one NaN value. If every row with an NaN was dropped, this would be far too much data lost to capture any useful analysis. Additionally, since I am looking at features individually instead of collectively, there is still a lot of useful data to look at, even if one or two pieces are missing from a row. Instead, I drop just the rows that have missing game titles, since that information is impossible to extrapolate.

Additionally, platform/genre/developer/rating all have to be translated into numeric values for some of the analysis.

Research Question(s)

The question I am aiming to answer is what markets a large video game developer should target when they develop new games in order to sell the most units. This includes genre, platform and geographic region. I first look at the entire dataset from 1980-2016, and then focus in on just to recent generations (PS3, PS4, Wii, WiiU, Xbox 360, Xbox One). It's important to note that since this dataset came out in 2016, the Nintendo Switch had not been released yet.

Methods


I primarily used data visualization when analyzing this data. Since I am looking at very specific and discrete categories (e.g. sales of Role-Playing games in Japan), bar charts were very useful in accurately depicting and comparing this information.

I start the analysis with a heatmap to see what categories might be correlated, which brings out some surprising information (developers and ratings are very closely correlated). Then I use a series of bar charts to look at game units sold in each geographic region based on genre, platform, developer.

I first look at the aggregate data from 1980-2016, and then focus in on just the recent generations of consoles.


Findings - Correlations

A heatmap showing correlation between game platform, genre, sales and scores. One would assume sales closely correlate between geographic regions, but it was surprising to see the strong correlation between game developer and game rating.



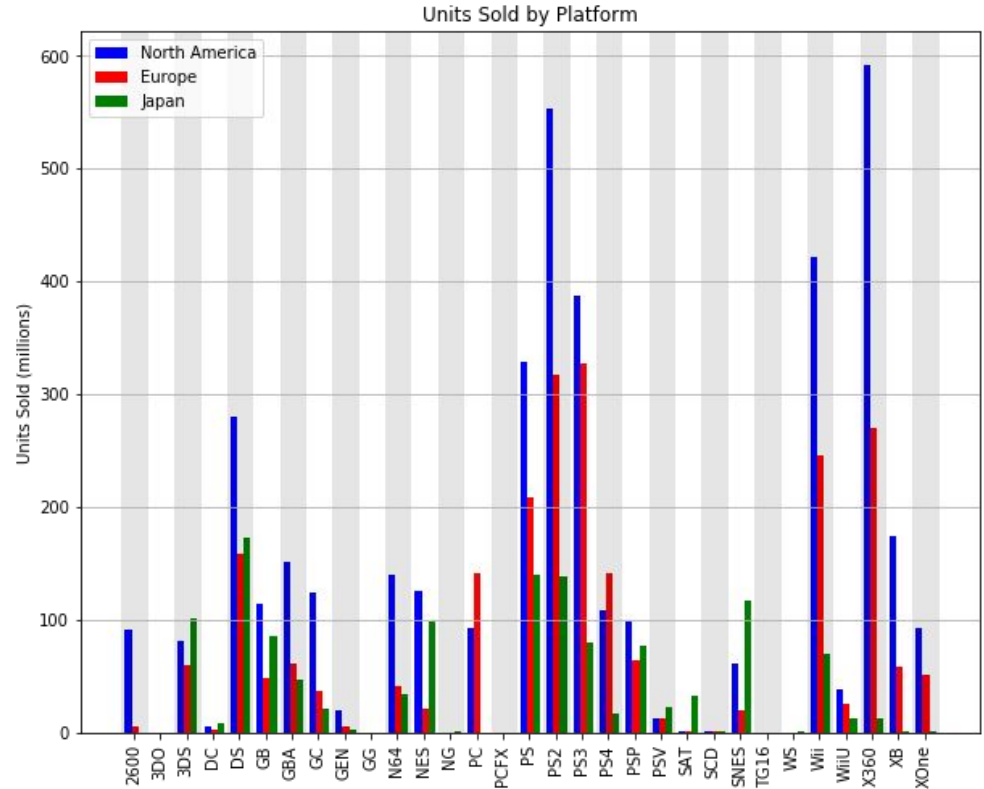
Findings - Total Units Sold by Genre

This shows total units sold by genre for the entire 38 year period. Interesting to note the difference in preferred genres based on geographic region. For instance, North America clearly likes Action games the most, while Japan enjoys Role-Playing games twice as much as any other genre.




Findings - Total Units Sold by Platform

This shows total game units sold for every game platform since the Atari 2600. You can tell pretty quickly that the PlayStation and Xbox generations have the most game sales, suggesting video games have gotten significantly more popular with each new console generation.



Findings - Total Units Sold by Developer

This shows total game units sold for the top 20 developers for the entire time frame. Nintendo primarily develops many of the games for their own systems, which is why they have so many more sales than any other developer.




Findings - Total Units Sold by Rating

This plot is different than previous plots in that we are no longer looking at geographic location, and instead focusing on game rating. If you are unfamiliar with the rating system, in order of increasing age:


- E: Everyone
- E10+: Everyone 10+
- T: Teen
- M: Mature

Since Nintendo is considered the “family-friendly” company, they distort this plot somewhat. I remove them in the next plot.



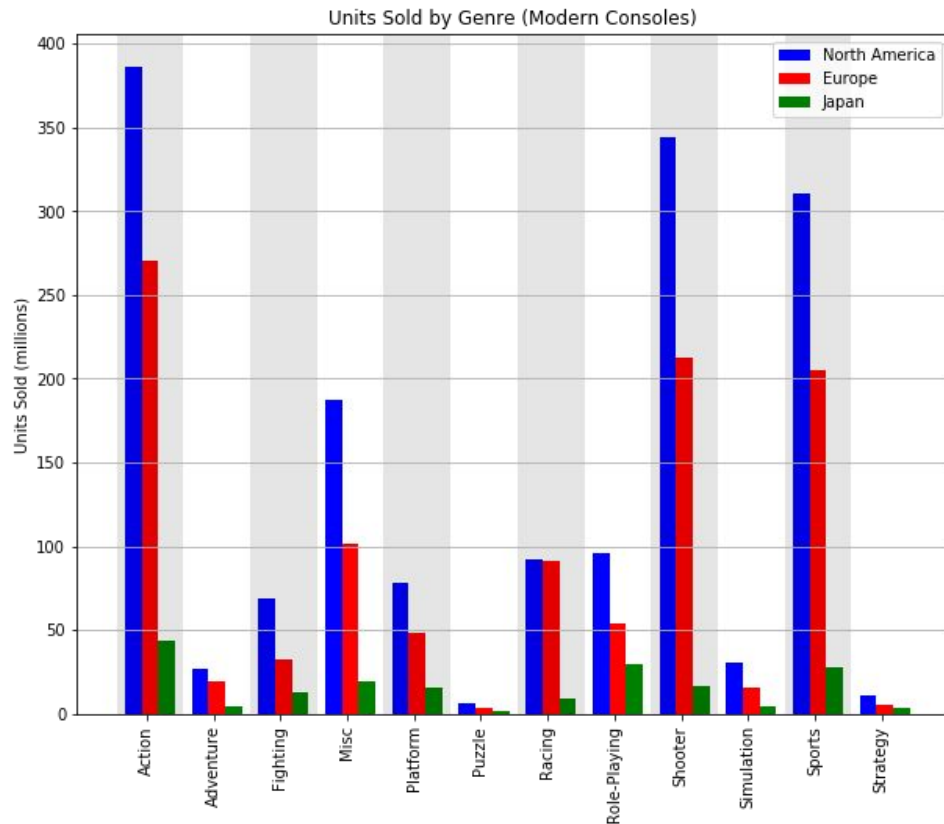
Findings - Total Units Sold by Rating

This shows the same information as the previous plot, with Nintendo removed. This allows us to have much better resolution and actually see how many units other game developers have sold. It's interesting to note how most of the companies skew towards one rating, e.g. EA Sports makes mostly rated-E games while Ubisoft Montreal makes mostly rated-M games. This makes sense since EA Sports focuses on sports-related games, while Ubisoft Montreal develops the Tom Clancy and Assassin's Creed games.



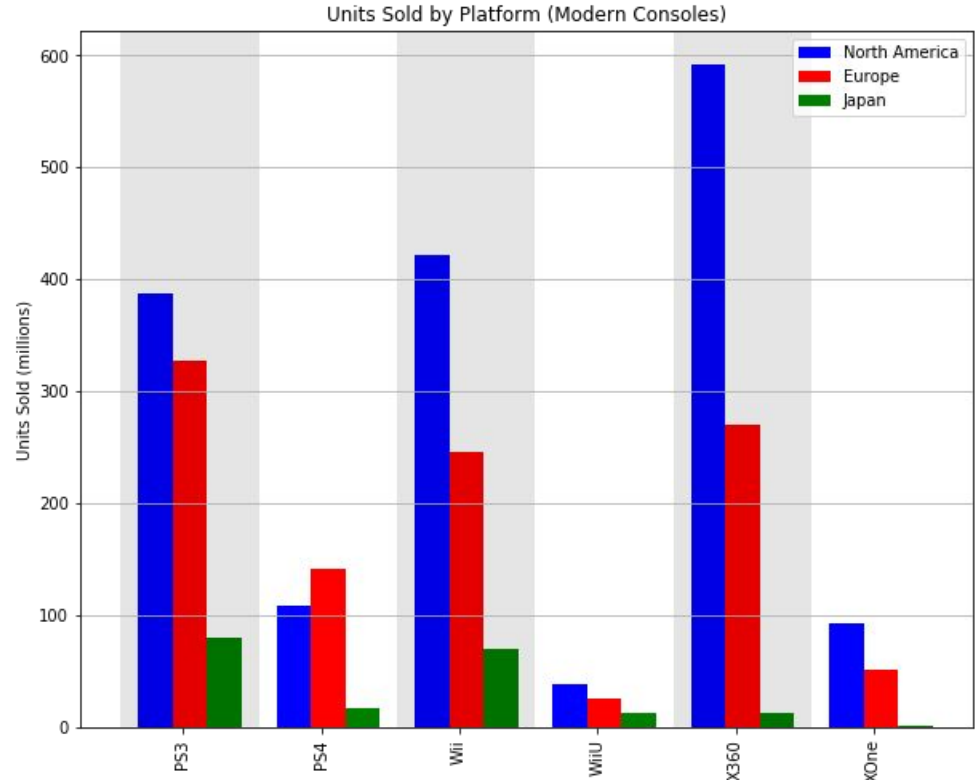
Findings - Modern Units Sold by Genre

This plot and the next plot focus in on the modern generations of consoles. This includes PS3, PS4, Wii, WiiU, Xbox 360 and Xbox One. Note that the dataset is from 2016, so is not completely up-to-date, and does not include the Nintendo Switch. Still, you can clearly see the North American and European preference for Action, Shooter and Sports games.



Findings - Total Units Sold by Platform

Finally, the total game units sold for each of the modern game consoles. Note that Xbox 360 clearly won the previous console generation, but it looks like PS4 is outpacing the Xbox One this generation. This shows that it would be beneficial for a developer to put out a game on multiple platforms in order to optimize sales.



Limitations

The dataset has one glaring limitation, which is that it is only relevant through 2016. Each year, the video game industry is more and more successful, so missing 2-3 years of recent data is significant. Also, this completely misses anything related to the Nintendo Switch, which has already had a massive impact on the industry.

Additionally, the conclusions drawn are only applicable to major video game developers who are seeking to maximize profit and game sales. This would not be applicable to independent game developers who are potentially more interested in unique games.

Conclusions

For the big game development studios, they should focus on Action, Shooter and Sports games in North America and Europe. They should ideally develop them on multiple platforms, since the Xbox 360 won the last console generation, while the PS4 looks to be winning this generation. In Japan, Role-Playing and Action games are roughly equal in appeal, with the PlayStation consoles (PS3 and PS4) being significantly more attractive than other consoles.

Acknowledgements

Dataset: <https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>

I did not discuss this particular project with anyone specifically, but I would like to thank my friends and acquaintances over many years who I have discussed these very topics with.

References

I did not use any references for this project - this was all completed on my own.

Analyzing Video Game Sales

February 18, 2019

1 Analyzing Video Game Sales

Jacob Miller The following dataset comes from Kaggle (<https://www.kaggle.com/rush4ratio/video-game-sales-with-ratings>) and captures video game sales and Metacritic ratings for video games ranging from 1980 to 2016. The data is organized by descending order of Global_Sales. I am going to look at sales by platform, genre, geographical region, developer and game rating. The goal will be to get an understanding of optimal markets to target as a video game developer.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('Video_Games_Sales_as_of_22_Dec_2016.csv')
```

Let's start by getting a basic understanding of what the data looks like. I'll look at the shape of the dataframe, the columns, the first 5 entries, and then focus in on the very first entry for clarity.

```
In [3]: df.shape
```

```
Out[3]: (16719, 16)
```

```
In [4]: df.columns
```

```
Out[4]: Index(['Name', 'Platform', 'Year_of_Release', 'Genre', 'Publisher', 'NA_Sales',
              'EU_Sales', 'JP_Sales', 'Other_Sales', 'Global_Sales', 'Critic_Score',
              'Critic_Count', 'User_Score', 'User_Count', 'Developer', 'Rating'],
              dtype='object')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	Name	Platform	Year_of_Release	Genre	Publisher
0	Wii Sports	Wii	2006.0	Sports	Nintendo
1	Super Mario Bros.	NES	1985.0	Platform	Nintendo
2	Mario Kart Wii	Wii	2008.0	Racing	Nintendo
3	Wii Sports Resort	Wii	2009.0	Sports	Nintendo
4	Pokemon Red/Pokemon Blue	GB	1996.0	Role-Playing	Nintendo

	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score \
0	41.36	28.96	3.77	8.45	82.53	76.0
1	29.08	3.58	6.81	0.77	40.24	NaN
2	15.68	12.76	3.79	3.29	35.52	82.0
3	15.61	10.93	3.28	2.95	32.77	80.0
4	11.27	8.89	10.22	1.00	31.37	NaN

	Critic_Count	User_Score	User_Count	Developer	Rating
0	51.0	8	322.0	Nintendo	E
1	NaN	NaN	NaN	NaN	NaN
2	73.0	8.3	709.0	Nintendo	E
3	73.0	8	192.0	Nintendo	E
4	NaN	NaN	NaN	NaN	NaN

```
In [6]: df.iloc[0]
```

```
Out[6]: Name           Wii Sports
Platform           Wii
Year_of_Release    2006
Genre             Sports
Publisher          Nintendo
NA_Sales           41.36
EU_Sales           28.96
JP_Sales           3.77
Other_Sales         8.45
Global_Sales       82.53
Critic_Score       76
Critic_Count       51
User_Score          8
User_Count         322
Developer          Nintendo
Rating             E
Name: 0, dtype: object
```

Now that we have a better understanding of the dataset, let's check what NaN values we have.

```
In [7]: df[df.isnull().any(axis = 1)].shape[0]
```

```
Out[7]: 9894
```

Ok, well that's a lot of rows to drop. Let's look at just the rows with NaN in the Name column. Even if one of the other columns is NaN (e.g. Critic_Score) we can still get useful information from most of the other columns.

```
In [8]: df[df['Name'].isnull()]
```

```
Out[8]:
```

	Name	Platform	Year_of_Release	Genre	Publisher	NA_Sales \
659	NaN	GEN	1993.0	NaN	Acclaim Entertainment	1.78
14246	NaN	GEN	1993.0	NaN	Acclaim Entertainment	0.00

	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	\
659	0.53	0.00	0.08	2.39	NaN	
14246	0.00	0.03	0.00	0.03	NaN	

	Critic_Count	User_Score	User_Count	Developer	Rating
659	NaN	NaN	NaN	NaN	NaN
14246	NaN	NaN	NaN	NaN	NaN

That's a lot more reasonable. We only have 2 games that are lacking names all together, so they really don't provide much useful information. Let's just drop these 2 rows. You can see from the shape afterwards that 2 rows were dropped.

```
In [9]: df.dropna(axis = 0, subset = ['Name'], inplace = True)
df.shape
```

```
Out[9]: (16717, 16)
```

We now have a really good understanding of how our data looks and how it is presented, so let's take a little detour and check out some of my favorite games. First, I'll look at how many Final Fantasy games are included in the entire dataset.

```
In [10]: df[df['Name'].str.contains('Final Fantasy')].shape[0]
```

```
Out[10]: 91
```

91 games is a little too many to display here, so I'll look at the top 25. Then, I'll focus in on my favorite one of the series, Final Fantasy IX.

```
In [11]: df[df['Name'].str.contains('Final Fantasy')][:25]
```

```
Out[11]:
```

	Name	Platform	Year_of_Release	\
65	Final Fantasy VII	PS	1997.0	
84	Final Fantasy X	PS2	2001.0	
88	Final Fantasy VIII	PS	1999.0	
148	Final Fantasy XII	PS2	2006.0	
173	Final Fantasy XIII	PS3	2009.0	
175	Final Fantasy IX	PS	2000.0	
177	Final Fantasy X-2	PS2	2003.0	
388	Final Fantasy III	SNES	1994.0	
431	Crisis Core: Final Fantasy VII	PSP	2007.0	
578	Final Fantasy XIII-2	PS3	2011.0	
632	Final Fantasy V	SNES	1992.0	
633	Final Fantasy Tactics	PS	1997.0	
723	Dissidia: Final Fantasy	PSP	2008.0	
759	Final Fantasy XIII	X360	2010.0	
776	Final Fantasy Tactics Advance	GBA	2003.0	
802	Final Fantasy III	DS	2006.0	
985	Final Fantasy II	SNES	1991.0	

1236	Final Fantasy XIV: A Realm Reborn	PC	2010.0
1267	Final Fantasy: Crystal Chronicles	GC	2003.0
1276	Dirge of Cerberus: Final Fantasy VII	PS2	2006.0
1383	Final Fantasy III	NES	1990.0
1411	Final Fantasy XII: Revenant Wings	DS	2007.0
1569	Final Fantasy X / X-2 HD Remaster	PS3	2013.0
1654	Final Fantasy IV	DS	2007.0
1703	Final Fantasy I & II: Dawn of Souls	GBA	2004.0

	Genre	Publisher	NA_Sales	EU_Sales	JP_Sales	\
65	Role-Playing	Sony Computer Entertainment	3.01	2.47	3.28	
84	Role-Playing	Sony Computer Entertainment	2.91	2.07	2.73	
88	Role-Playing	SquareSoft	2.28	1.72	3.63	
148	Role-Playing	Square Enix	1.88	0.00	2.33	
173	Role-Playing	Square Enix	1.74	1.21	1.87	
175	Role-Playing	SquareSoft	1.62	0.77	2.78	
177	Role-Playing	Electronic Arts	1.92	1.08	2.11	
388	Role-Playing	SquareSoft	0.86	0.00	2.55	
431	Role-Playing	Square Enix	1.35	0.59	0.80	
578	Role-Playing	Square Enix	0.78	0.73	0.89	
632	Role-Playing	SquareSoft	0.00	0.00	2.43	
633	Role-Playing	SquareSoft	0.93	0.12	1.34	
723	Fighting	Square Enix	0.51	0.50	0.91	
759	Role-Playing	Square Enix	1.28	0.67	0.01	
776	Role-Playing	SquareSoft	0.82	0.37	0.89	
802	Role-Playing	Square Enix	0.89	0.04	1.07	
985	Role-Playing	Square	0.24	0.09	1.33	
1236	Role-Playing	Square Enix	0.88	0.48	0.00	
1267	Role-Playing	Nintendo	0.72	0.38	0.36	
1276	Shooter	Square Enix	0.47	0.37	0.52	
1383	Role-Playing	SquareSoft	0.00	0.00	1.39	
1411	Role-Playing	Square Enix	0.33	0.41	0.54	
1569	Role-Playing	Square Enix	0.43	0.36	0.32	
1654	Simulation	Square Enix	0.51	0.04	0.62	
1703	Role-Playing	Nintendo	0.64	0.24	0.29	

	Other_Sales	Global_Sales	Critic_Score	Critic_Count	User_Score	\
65	0.96	9.72	92.0	20.0	9.2	
84	0.33	8.05	92.0	53.0	8.7	
88	0.23	7.86	90.0	24.0	8.6	
148	1.74	5.95	92.0	64.0	7.6	
173	0.51	5.33	83.0	83.0	7.3	
175	0.14	5.30	94.0	22.0	8.9	
177	0.17	5.29	85.0	45.0	6.6	
388	0.02	3.42	NaN	NaN	NaN	
431	0.43	3.18	83.0	67.0	8	
578	0.23	2.63	79.0	53.0	6.6	
632	0.02	2.45	NaN	NaN	NaN	

633	0.06	2.45	83.0	12.0	8.2
723	0.31	2.23	79.0	61.0	8
759	0.21	2.16	82.0	54.0	6.3
776	0.05	2.13	87.0	44.0	7.8
802	0.09	2.08	77.0	45.0	7.1
985	0.12	1.77	NaN	NaN	NaN
1236	0.17	1.52	NaN	NaN	NaN
1267	0.04	1.49	80.0	55.0	8.1
1276	0.12	1.48	57.0	51.0	6.9
1383	0.01	1.40	NaN	NaN	NaN
1411	0.10	1.37	81.0	44.0	6.1
1569	0.16	1.27	85.0	50.0	8.5
1654	0.05	1.21	85.0	52.0	7.7
1703	0.02	1.19	NaN	NaN	NaN

	User_Count	Developer	Rating
65	1282.0	SquareSoft	T
84	1056.0	SquareSoft	T
88	644.0	SquareSoft	T
148	972.0	Square Enix	T
173	2483.0	Square Enix	T
175	779.0	SquareSoft	T
177	400.0	SquareSoft	T
388	NaN	NaN	NaN
431	463.0	Square Enix	T
578	707.0	Square Enix	T
632	NaN	NaN	NaN
633	190.0	SquareSoft	T
723	139.0	Square Enix	T
759	539.0	Square Enix	T
776	212.0	Square Enix	E
802	136.0	Matrix Software	E10+
985	NaN	NaN	NaN
1236	NaN	NaN	NaN
1267	94.0	Square Enix	T
1276	103.0	Square Enix	T
1383	NaN	NaN	NaN
1411	60.0	Square Enix, Think and Feel	E10+
1569	332.0	Virtuos	T
1654	164.0	Matrix Software	E10+
1703	NaN	NaN	NaN

```
In [12]: df[df['Name'] == 'Final Fantasy IX'].T
```

```
Out[12]:
```

Name	Final Fantasy IX
Platform	PS
Year_of_Release	2000

Genre	Role-Playing
Publisher	SquareSoft
NA_Sales	1.62
EU_Sales	0.77
JP_Sales	2.78
Other_Sales	0.14
Global_Sales	5.3
Critic_Score	94
Critic_Count	22
User_Score	8.9
User_Count	779
Developer	SquareSoft
Rating	T

Next, let's start doing a little analysis on the dataset. In order to not skew the analysis based on a few outlier reviews, I'm going to drop any rows that both less than 5 Critic_Scores and less than 5 User_Scores. Below, you can see that User_Score is a string, so we have to cast it as a number to be able to check how many reviews there actually are.

```
In [13]: print('Critic_Score type: ' + str(type(df['Critic_Score'][0])))
        print('User_Score type: ' + str(type(df['User_Score'][0])))
```

```
Critic_Score type: <class 'numpy.float64'>
User_Score type: <class 'str'>
```

```
In [14]: try:
        df['User_Score'] = df['User_Score'].astype(np.float64)
    except ValueError as e:
        print('--- Oops! --- \nValueError: ' + str(e))
```

```
--- Oops! ---
ValueError: could not convert string to float: 'tbd'
```

Huh... ok. So numbers within the User_Score column are stored as strings, but there are *also* non-numeric string values. Let's see how many of those we have.

```
In [15]: len(df[df['User_Score'] == 'tbd'])
```

```
Out[15]: 2425
```

Wow. That's a lot of scores to be determined, so now let's drop them all. If we are doing an investigation of ratings, it won't help us much to carry around a bunch of useless ratings.

```
In [16]: df.drop(df[(df['User_Score'] == 'tbd')].index, inplace = True)
        len(df[df['User_Score'] == 'tbd'])
```

```
Out[16]: 0
```

Now that we've cleared that up, we'll cast the User_Scores into a numeric-type, and then we'll look at which games have less than 5 reviews.

```
In [17]: df['User_Score'] = df['User_Score'].astype(np.float64)
         df[(df['Critic_Count'] < 5) & (df['User_Count'] < 5)]
```

```
Out[17]:
```

	Name	Platform	Year_of_Release	\
1567	Battle Arena Toshinden	PS	1994.0	
2420	Ford Racing	PS	2001.0	
3924	The BIGS	PS2	2007.0	
4219	Dragon Ball Z: Collectible Card Game	GBA	2002.0	
4689	Marvel Super Hero Squad	DS	2009.0	
5029	Get Fit with Mel B	PS3	2010.0	
6994	Gundam Battle Assault	PS	1998.0	
7367	SBK X: Superbike World Championship	PS3	2010.0	
8075	Six Flags Fun Park	DS	2008.0	
8701	Chaotic: Shadow Warriors	X360	2009.0	
9774	NCIS	PS3	2011.0	
10157	Cloudy With a Chance of Meatballs	Wii	2009.0	
10177	Astro Boy: The Video Game	PS2	2009.0	
10618	Sudoku Mania	DS	2006.0	
10656	Rolling Stone: Drum King	Wii	2009.0	
11258	Puss in Boots	PS3	2011.0	
11549	Hidden Invasion	PS2	2001.0	
13413	Cabela's Trophy Bucks	X360	2007.0	
13632	LEGO Friends	3DS	2013.0	
14801	MX vs. ATV Supercross	PS3	2014.0	
14891	Madden NFL 07	GBA	2006.0	
14897	AniMates!	DS	2007.0	
15044	NBA Starting Five	XB	2002.0	

	Genre	Publisher	NA_Sales	EU_Sales	\
1567	Fighting	Sony Computer Entertainment	0.39	0.26	
2420	Racing	Empire Interactive	0.48	0.33	
3924	Sports	Take-Two Interactive	0.25	0.19	
4219	Misc	Infogrames	0.33	0.12	
4689	Fighting	THQ	0.36	0.02	
5029	Sports	Black Bean Games	0.15	0.16	
6994	Fighting	Namco Bandai Games	0.13	0.09	
7367	Racing	Black Bean Games	0.05	0.12	
8075	Misc	Brash Entertainment	0.17	0.00	
8701	Action	Activision	0.14	0.00	
9774	Adventure	Ubisoft	0.07	0.04	
10157	Platform	Ubisoft	0.10	0.00	
10177	Action	D3Publisher	0.05	0.04	
10618	Puzzle	Zoo Digital Publishing	0.09	0.00	
10656	Misc	505 Games	0.09	0.00	
11258	Action	THQ	0.07	0.01	

11549	Shooter	Swing! Entertainment	0.04	0.03
13413	Sports	Activision Value	0.04	0.00
13632	Action	Warner Bros. Interactive Entertainment	0.00	0.04
14801	Racing	Nordic Games	0.02	0.01
14891	Sports	Electronic Arts	0.02	0.01
14897	Simulation	Lexicon Entertainment	0.02	0.00
15044	Sports	Konami Digital Entertainment	0.02	0.01

	JP_Sales	Other_Sales	Global_Sales	Critic_Score	Critic_Count	\
1567	0.53	0.08	1.27	69.0	4.0	
2420	0.00	0.06	0.86	53.0	4.0	
3924	0.00	0.06	0.51	71.0	4.0	
4219	0.00	0.01	0.46	62.0	4.0	
4689	0.00	0.03	0.41	61.0	4.0	
5029	0.00	0.07	0.38	73.0	4.0	
6994	0.00	0.02	0.23	61.0	4.0	
7367	0.00	0.04	0.21	73.0	4.0	
8075	0.00	0.01	0.18	34.0	4.0	
8701	0.00	0.01	0.16	60.0	4.0	
9774	0.00	0.02	0.12	50.0	4.0	
10157	0.00	0.01	0.11	68.0	4.0	
10177	0.00	0.01	0.11	56.0	4.0	
10618	0.00	0.01	0.10	25.0	4.0	
10656	0.00	0.01	0.10	32.0	4.0	
11258	0.00	0.01	0.09	70.0	4.0	
11549	0.00	0.01	0.08	39.0	4.0	
13413	0.00	0.00	0.05	39.0	4.0	
13632	0.00	0.00	0.04	43.0	4.0	
14801	0.00	0.00	0.03	58.0	4.0	
14891	0.00	0.00	0.03	68.0	4.0	
14897	0.00	0.00	0.03	33.0	4.0	
15044	0.00	0.00	0.02	48.0	4.0	

	User_Score	User_Count	Developer	Rating
1567	6.3	4.0	Tamsoft	T
2420	5.8	4.0	Toolbox Design	E
3924	4.8	4.0	Blue Castle Games	E
4219	5.5	4.0	ImaginEngine	E
4689	2.8	4.0	THQ	E10+
5029	5.8	4.0	Lightning Fish Games	E
6994	7.8	4.0	Natsume	T
7367	7.5	4.0	Milestone S.r.l	E10+
8075	7.0	4.0	Brash Entertainment	E
8701	5.0	4.0	FUN Labs	E10+
9774	4.0	4.0	Ubisoft	T
10157	5.3	4.0	Ubisoft Shanghai	E
10177	5.0	4.0	High Voltage Software	E10+
10618	3.5	4.0	FrontLine Studios	E

10656	1.8	4.0	505 Games	T
11258	7.0	4.0	THQ	E10+
11549	5.8	4.0	Toka	T
13413	3.0	4.0	FUN Labs	T
13632	8.3	4.0	Hellbent Games	E
14801	3.5	4.0	Nordic Games Publishing	E
14891	9.3	4.0	Exient Entertainment	E
14897	2.5	4.0	Dreamcatcher	E
15044	5.5	4.0	Konami	E

Great. So as badly as someone may want to "Get Fit With Mel B", apparently not many other people did. Let's take a look at the shape of the dataframe before dropping these games, then we'll drop them and take a look at the shape afterwards to make sure everything went ok. We can tell that 23 games were dropped.

```
In [18]: print('--- Shape before dropping games with <5 reviews ---\n' + \
            str(df.shape))

df.drop(df[(df['Critic_Count'] < 5) & (df['User_Count'] < 5)].index,
        inplace = True)

print('\n--- Shape after dropping games with <5 reviews ---\n' + \
      str(df.shape))
```

```
--- Shape before dropping games with <5 reviews ---
(14292, 16)
```

```
--- Shape after dropping games with <5 reviews ---
(14269, 16)
```

I'd like to see what the correlation is between certain categories in this dataframe, but in order to do that, we need to get the categories we want to look at into numeric values. We'll start by creating a new "integer_df"

```
In [19]: integer_df = df.copy()
```

This is where things get a little backwards. The goal is to convert all instances of one string value (e.g. "E10+") into a numeric representation (e.g. "1"). However, it turns out some of the Developers and Ratings are actually being interpreted as numeric values, even though they are obviously strings (e.g. "E10+"). For example, you can see below that the `np.unique` is trying (and failing) to compare floats and strings. I won't explain in detail how I was able to get to the bottom of this, but in order to resolve this, first we have to cast **all** values into strings, and then convert them to a numeric representation.

```
In [20]: try:
            np.unique(df['Rating'])
        except TypeError as e:
            print('--- Oops! --- \nTypeError: ' + str(e))
```

--- Oops! ---

TypeError: '<' not supported between instances of 'float' and 'str'

```
In [21]: integer_df['Developer'] = integer_df['Developer'].astype(str)
integer_df['Rating'] = integer_df['Rating'].astype(str)
```

```
In [22]: def integer_dictionary(string_column):
    """
    Returns a dictionary of {string : integer} for a column/series in dataframe
    """
    return {string : i for i, string in \
            enumerate(np.unique(integer_df[string_column]))}

platform_ints = integer_dictionary('Platform')
genre_ints = integer_dictionary('Genre')
developer_ints = integer_dictionary('Developer')
rating_ints = integer_dictionary('Rating')
```

```
In [23]: def apply_int_trans(string_column, integer_dict):
    """
    Translates all instances of a string to an integer within one column
    """
    return integer_df[string_column].apply(lambda x: integer_dict[x])

integer_df['Platform'] = apply_int_trans('Platform', platform_ints)
integer_df['Genre'] = apply_int_trans('Genre', genre_ints)
integer_df['Developer'] = apply_int_trans('Developer', developer_ints)
integer_df['Rating'] = apply_int_trans('Rating', rating_ints)
```

You can see below that Platform, Genre, Developer and Rating have now all been converted to numeric representations. We can now use these to plot a heatmap and see what type of correlations we have.

```
In [24]: integer_df.head()
```

```
Out[24]:
```


	Name	Platform	Year_of_Release	Genre	Publisher	\
0	Wii Sports	26	2006.0	10	Nintendo	
1	Super Mario Bros.	11	1985.0	4	Nintendo	
2	Mario Kart Wii	26	2008.0	6	Nintendo	
3	Wii Sports Resort	26	2009.0	10	Nintendo	
4	Pokemon Red/Pokemon Blue	5	1996.0	7	Nintendo	

	NA_Sales	EU_Sales	JP_Sales	Other_Sales	Global_Sales	Critic_Score	\
0	41.36	28.96	3.77	8.45	82.53	76.0	
1	29.08	3.58	6.81	0.77	40.24	NaN	
2	15.68	12.76	3.79	3.29	35.52	82.0	
3	15.61	10.93	3.28	2.95	32.77	80.0	
4	11.27	8.89	10.22	1.00	31.37	NaN	

	Critic_Count	User_Score	User_Count	Developer	Rating
0	51.0	8.0	322.0	835	1
1	NaN	NaN	NaN	1380	8
2	73.0	8.3	709.0	835	1
3	73.0	8.0	192.0	835	1
4	NaN	NaN	NaN	1380	8

```
In [25]: correlation_features = ['Platform',
                                'Genre',
                                'NA_Sales',
                                'EU_Sales',
                                'JP_Sales',
                                'Global_Sales',
                                'Critic_Score',
                                'User_Score',
                                'Developer',
                                'Rating']

sns.heatmap(integer_df[correlation_features].corr(), cmap = 'YlGnBu')
plt.title('Correlation between Platform, Genre, \nSales and Game Scores')
plt.show()
```



Heatmaps are really fun and can help tease out some correlations in the data. For instance, it's probably pretty obvious that sales between various geographic regions are correlated, but I wouldn't have thought that Ratings and Developers were so closely correlated.

Let's look at how game sales break down by genre and by platform. If I was a game developer, this information would definitely help me decide what kind of game and platform to target. One thing that pops out right away is the difference between the Japanese and North American sales in the Role-Playing and Shooter genres. **Role-Playing** games have much higher sales than any other genre in Japan, whereas Shooters are actually the least sold Genre. On the other hand, in North America, **Shooters** are the third most-sold genre and Role-Playing is the seventh.

```
In [26]: def sales_by_genre(dataframe, region):
        '''
        Returns a dictionary of {genre : sales in millions of units} for "region"
        '''
        return {genre : dataframe[dataframe['Genre'] == genre][region].sum() \
                for genre in genre_ints.keys()}

na_sales_by_genre = sales_by_genre(df, 'NA_Sales')
eu_sales_by_genre = sales_by_genre(df, 'EU_Sales')
jp_sales_by_genre = sales_by_genre(df, 'JP_Sales')

genre_indices = range(len(na_sales_by_genre))
genre_width = np.min(np.diff(genre_indices))/4


fig, ax = plt.subplots(figsize = (10, 8))
bar1 = ax.bar(genre_indices - genre_width, na_sales_by_genre.values(),
              genre_width, color = 'blue')
bar2 = ax.bar(genre_indices, eu_sales_by_genre.values(),
              genre_width, color = 'red')
bar3 = ax.bar(genre_indices + genre_width, jp_sales_by_genre.values(),
              genre_width, color = 'green')

plt.xticks(rotation = 90)
ax.set_xticks(genre_indices - genre_width/5)
ax.set_xticklabels((na_sales_by_genre.keys()))

plt.title('Units Sold by Genre')
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['North America', 'Europe', 'Japan'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_genre), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()
```



```
In [27]: def sales_by_platform(region):
        '''
        Returns a dictionary of {platform : sales in millions of units}
        for "region"
        '''
        return {platform : df[df['Platform'] == platform][region].sum() for \
                platform in platform_ints.keys()}

na_sales_by_platform = sales_by_platform('NA_Sales')
eu_sales_by_platform = sales_by_platform('EU_Sales')
jp_sales_by_platform = sales_by_platform('JP_Sales')

platform_indices = range(len(na_sales_by_platform))
platform_width = np.min(np.diff(platform_indices))/4

fig, ax = plt.subplots(figsize = (10, 8))
```



```

bar1 = ax.bar(platform_indices - platform_width, na_sales_by_platform.values(),
               platform_width, color = 'blue')
bar2 = ax.bar(platform_indices, eu_sales_by_platform.values(),
               platform_width, color = 'red')
bar3 = ax.bar(platform_indices + platform_width, jp_sales_by_platform.values(),
               platform_width, color = 'green')

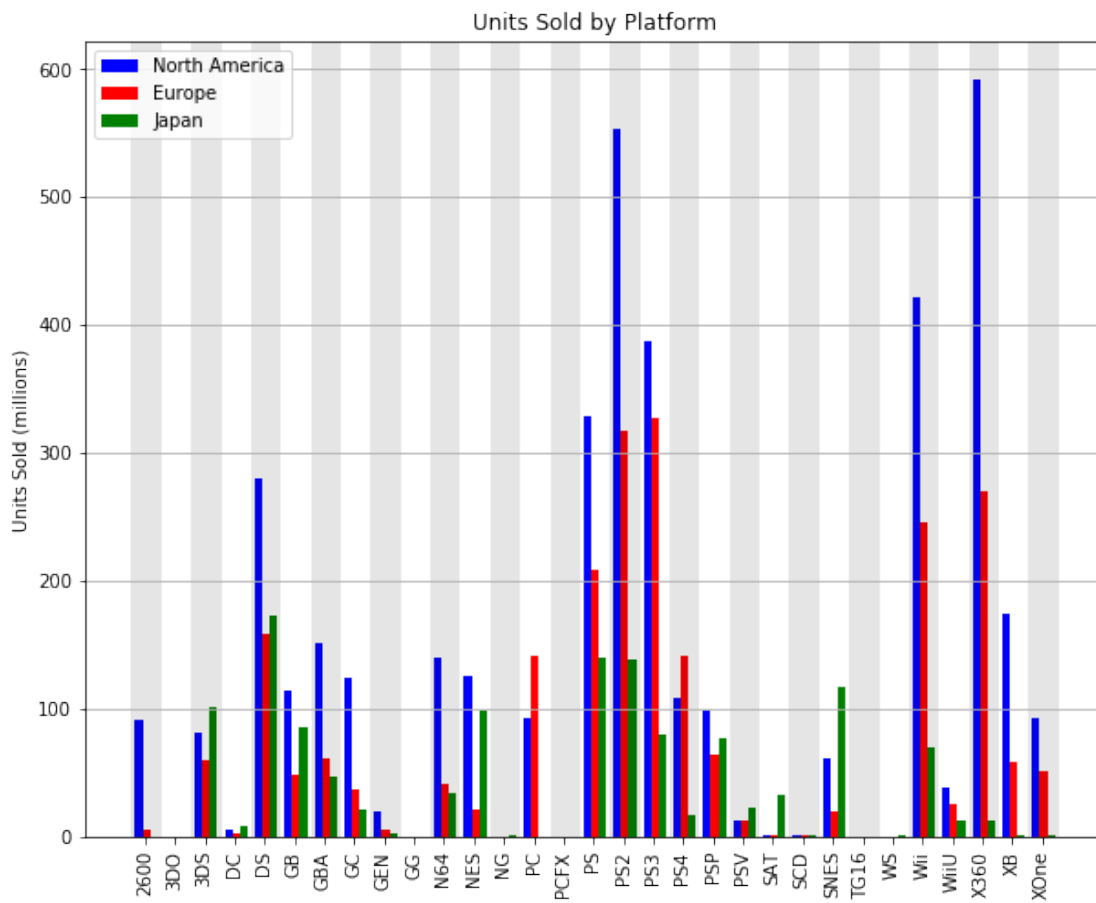
plt.xticks(rotation = 90)
ax.set_xticks(platform_indices - platform_width/5)
ax.set_xticklabels((na_sales_by_platform.keys()))

plt.title('Units Sold by Platform')
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['North America', 'Europe', 'Japan'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_platform), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()

```



In order to figure out what the top developers are (by number of games developed), we'll get all of the unique developers and the number of times they show up. Then, we'll create a dictionary that will allow us to first sort the developers by number of games developed, and then subsequently lookup developers numerically.

```
In [28]: devs, dev_counts = np.unique(integer_df['Developer'], return_counts = True)
        dev_counts_dict = {i : [dev_counts, i] for i, dev_counts in \
                               enumerate(dev_counts)}
        top_devs_integers = sorted(dev_counts_dict.values(), reverse = True)

        rev = dict({(v, k) for k, v in developer_ints.items()})
```

```
In [29]: rev[1000]
```

```
Out[29]: 'Rising Star Games'
```

Great! Now let's see who our top developer is...

```
In [30]: rev[top_devs_integers[0][1]]
```

```
Out[30]: 'nan'
```

Oops... Remember, we didn't want to drop 9000+ rows of data, but that means some NaNs will show up in our data. That's fine - we can work around them when we know where they are. So let's just look at the top 20 developers that are **not** NaN.

```
In [31]: top_20_integers = top_devs_integers[1:21]
        top_20_devs = [rev[dev[1]] for dev in top_20_integers]
        print('\n--- Top 20 Developers ---')
        for i, dev in enumerate(top_20_devs):
            print(str(i + 1) + '. ' + dev)
```

```
--- Top 20 Developers ---
```

1. EA Sports
2. EA Canada
3. Capcom
4. Ubisoft
5. Konami
6. EA Tiburon
7. Ubisoft Montreal
8. Visual Concepts
9. Omega Force
10. Electronic Arts
11. TT Games
12. Traveller's Tales
13. Nintendo
14. Vicarious Visions

15. Codemasters
16. Yuke's
17. Namco
18. Maxis
19. Neversoft Entertainment
20. Midway

Some familiar names there, but I'd never heard of "Vicarious Visions" before. Let's briefly investigate them as an aside. Turns out they developed many of the Guitar Hero and Crash Bandicoot games. Neat!

```
In [32]: df[df['Developer'] == 'Vicarious Visions']\
        [['Name', 'Developer', 'Critic_Score']].head(20)
```

```
Out [32]:
```

	Name	Developer	Critic_Score
234	Guitar Hero III: Legends of Rock	Vicarious Visions	86.0
351	Guitar Hero: World Tour	Vicarious Visions	86.0
383	Guitar Hero: On Tour	Vicarious Visions	71.0
519	Finding Nemo	Vicarious Visions	NaN
873	Crash Nitro Kart	Vicarious Visions	69.0
1032	Crash Bandicoot: The Huge Adventure	Vicarious Visions	78.0
1181	Guitar Hero: On Tour Decades	Vicarious Visions	72.0
1205	Spider-Man 2: Enter: Electro	Vicarious Visions	74.0
1206	Guitar Hero 5	Vicarious Visions	89.0
1421	Crash Bandicoot 2: N-Tranced	Vicarious Visions	75.0
1454	Doom 3	Vicarious Visions	
1478	Skylanders SWAP Force	Vicarious Visions	
1547	Guitar Hero: Aerosmith	Vicarious Visions	
1625	SpongeBob SquarePants: Battle for Bikini Bottom	Vicarious Visions	
1692	SpongeBob SquarePants: Revenge of the Flying D...	Vicarious Visions	
1715	Tony Hawk's Pro Skater 2	Vicarious Visions	
2034	Skylanders SWAP Force	Vicarious Visions	
2128	Band Hero	Vicarious Visions	
2294	Marvel: Ultimate Alliance 2	Vicarious Visions	
2310	Spider-Man 3	Vicarious Visions	

1454	88.0
1478	83.0
1547	73.0
1625	NaN
1692	71.0
1715	95.0
2034	83.0
2128	79.0
2294	73.0
2310	50.0

We are now ready to look at sales based on developers. Doesn't come as much of a shock that Nintendo leads this category, since they primarily develop many of the games for their own systems.

```
In [33]: def sales_by_developer(region):
        '''
        Returns a dictionary of {developer : sales in millions of units}
        for "region"
        '''
        return {developer : df[df['Developer'] == developer][region].sum() for \
                developer in top_20_devs}

na_sales_by_developer = sales_by_developer('NA_Sales')
eu_sales_by_developer = sales_by_developer('EU_Sales')
jp_sales_by_developer = sales_by_developer('JP_Sales')

developer_indices = range(len(na_sales_by_developer))
developer_width = np.min(np.diff(developer_indices))/4

fig, ax = plt.subplots(figsize = (10, 8))

bar1 = ax.bar(developer_indices - developer_width,
              na_sales_by_developer.values(),
              developer_width, color = 'blue')
bar2 = ax.bar(developer_indices,
              eu_sales_by_developer.values(),
              developer_width, color = 'red')
bar3 = ax.bar(developer_indices + developer_width,
              jp_sales_by_developer.values(),
              developer_width, color = 'green')


plt.xticks(rotation = 90)
ax.set_xticks(developer_indices - developer_width/5)
ax.set_xticklabels((na_sales_by_developer.keys()))

plt.title('Units Sold for Top 20 Developers')
plt.ylabel('Units Sold (millions)')
```

```
plt.legend(labels = ['North America', 'Europe', 'Japan'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_developer), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()
```



Next, we are going to look at game sales by rating to get an idea of what type of rating would sell the most. Current ratings are (in ascending order of age group) E, E10+, T, M, AO. Below, you'll also see a couple of "K-A" which was what the "E" rating was known as prior to 1998, and "RP", or Rating Pending. Since these are such low counts, I am going to not carry them through (along with the single AO rated game).

```
In [34]: df['Rating'] = df['Rating'].astype(str)
ratings, ratings_counts = np.unique(df['Rating'], return_counts = True)
```

```

        for i in zip(ratings, ratings_counts):
            print(i)

('AO', 1)
('E', 2403)
('E10+', 1045)
('EC', 1)
('K-A', 3)
('M', 1521)
('RP', 3)
('T', 2575)
('nan', 6717)

In [35]: def sales_by_rating(rating, dev_list):
        '''
        Returns dictionary of {rating : sales in millions of units} for each
        "developer"
        '''
        return {developer : df[(df['Developer'] == developer) & \
                                (df['Rating'] == rating)][ 'Global_Sales'].sum() \
                                for developer in dev_list}

e_sales_by_developer = sales_by_rating('E', top_20_devs)
e10_sales_by_developer = sales_by_rating('E10+', top_20_devs)
t_sales_by_developer = sales_by_rating('T', top_20_devs)
m_sales_by_developer = sales_by_rating('M', top_20_devs)

rating_indices = range(len(e_sales_by_developer))
rating_width = np.min(np.diff(rating_indices))/5

fig, ax = plt.subplots(figsize = (10, 8))

bar1 = ax.bar(rating_indices - 1.5*rating_width, e_sales_by_developer.values(),
              rating_width, color = 'blue')
bar2 = ax.bar(rating_indices - 0.5*rating_width, e10_sales_by_developer.values(),
              rating_width, color = 'red')
bar3 = ax.bar(rating_indices + 0.5*rating_width, t_sales_by_developer.values(),
              rating_width, color = 'green')
bar4 = ax.bar(rating_indices + 1.5*rating_width, m_sales_by_developer.values(),
              rating_width, color = 'orange')

plt.xticks(rotation = 90)
ax.set_xticks(developer_indices - developer_width/5)
ax.set_xticklabels((na_sales_by_developer.keys()))


plt.title('Units Sold by Rating for Top 20 Developers')

```

```
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['E', 'E10+', 'T', 'M'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_developer), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()
```



Well that's a pretty telling graph: Nintendo makes a **LOT** of rated-E games. However, it makes it pretty difficult to see any resolution on the other developers, so let's briefly disregard Nintendo and look at this again.

```
In [36]: top_19_devs = [dev for dev in top_20_devs if not dev == 'Nintendo']
```

```

e_sales_by_developer = sales_by_rating('E', top_19_devs)
e10_sales_by_developer = sales_by_rating('E10+', top_19_devs)
t_sales_by_developer = sales_by_rating('T', top_19_devs)
m_sales_by_developer = sales_by_rating('M', top_19_devs)

rating_indices = range(len(e_sales_by_developer))
rating_width = np.min(np.diff(rating_indices))/5

fig, ax = plt.subplots(figsize = (10, 8))

bar1 = ax.bar(rating_indices - 1.5*rating_width, e_sales_by_developer.values(),
              rating_width, color = 'blue')
bar2 = ax.bar(rating_indices - 0.5*rating_width, e10_sales_by_developer.values(),
              rating_width, color = 'red')
bar3 = ax.bar(rating_indices + 0.5*rating_width, t_sales_by_developer.values(),
              rating_width, color = 'green')
bar4 = ax.bar(rating_indices + 1.5*rating_width, m_sales_by_developer.values(),
              rating_width, color = 'orange')


plt.xticks(rotation = 90)
ax.set_xticks(developer_indices - developer_width/5)
ax.set_xticklabels(top_19_devs)

plt.title('Units Sold by Rating for Top 19 Developers (Dropping Nintendo)')
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['E', 'E10+', 'T', 'M'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_developer), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()

```

This has been a pretty exhaustive exploration of game sales over the last almost-40 years. However, what if we were interested in developing a game for the recent generation of consoles? Keep in mind, this dataset came out in 2016, so it doesn't include the Nintendo Switch, but will still give us a good picture.

```
In [37]: modern = ['PS3', 'PS4', 'Wii', 'WiiU', 'X360', 'XOne']
         modern_df = pd.DataFrame()
         for platform in modern:
             modern_df = pd.concat([modern_df, df[df['Platform'] == platform]])
```

```
In [38]: modern_df.shape
```

```
Out[38]: (4058, 16)
```

```
In [39]: na_sales_by_genre = sales_by_genre(modern_df, 'NA_Sales')
         eu_sales_by_genre = sales_by_genre(modern_df, 'EU_Sales')
```

```

jp_sales_by_genre = sales_by_genre(modern_df, 'JP_Sales')

genre_indices = range(len(na_sales_by_genre))
genre_width = np.min(np.diff(genre_indices))/4

fig, ax = plt.subplots(figsize = (10, 8))
bar1 = ax.bar(genre_indices - genre_width, na_sales_by_genre.values(),
              genre_width, color = 'blue')
bar2 = ax.bar(genre_indices, eu_sales_by_genre.values(),
              genre_width, color = 'red')
bar3 = ax.bar(genre_indices + genre_width, jp_sales_by_genre.values(),
              genre_width, color = 'green')


plt.xticks(rotation = 90)
ax.set_xticks(genre_indices - genre_width/5)
ax.set_xticklabels((na_sales_by_genre.keys()))

plt.title('Units Sold by Genre (Modern Consoles)')
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['North America', 'Europe', 'Japan'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_genre), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()

```



```
In [40]: def sales_by_platform(region):
        '''
        Returns a dictionary of {platform : sales in millions of units}
        for "region"
        '''
        return {platform : modern_df[modern_df['Platform'] == platform][region].\
                sum() for platform in modern}

na_sales_by_platform = sales_by_platform('NA_Sales')
eu_sales_by_platform = sales_by_platform('EU_Sales')
jp_sales_by_platform = sales_by_platform('JP_Sales')

platform_indices = range(len(modern))
platform_width = np.min(np.diff(platform_indices))/4

fig, ax = plt.subplots(figsize = (10, 8))
```

```

bar1 = ax.bar(platform_indices - platform_width, na_sales_by_platform.values(),
               platform_width, color = 'blue')
bar2 = ax.bar(platform_indices, eu_sales_by_platform.values(),
               platform_width, color = 'red')
bar3 = ax.bar(platform_indices + platform_width, jp_sales_by_platform.values(),
               platform_width, color = 'green')


plt.xticks(rotation = 90)
ax.set_xticks(platform_indices - platform_width/5)
ax.set_xticklabels((na_sales_by_platform.keys()))

plt.title('Units Sold by Platform (Modern Consoles)')
plt.ylabel('Units Sold (millions)')
plt.legend(labels = ['North America', 'Europe', 'Japan'], loc = 'best')
plt.grid(which = 'major', axis = 'y')

for i in range(0, len(na_sales_by_platform), 2):
    plt.axvspan(i - 0.5, i + 0.5, facecolor = 'black', alpha = 0.1)

plt.show()

```



So there you have it - of the big-name developers, the ones having the most success (by far!) are selling **Action**, **Shooter**, or **Sports** games. Keeping in mind this dataset is at least 2 years old, there have been a lot of data generated for the current generation of consoles (PS4, Xbox One, Nintendo Switch) that is not captured here. Last generation, the **Xbox 360** was a clear winner. However, it's looking like the **PS4** is taking the crown so far in this generation!